


FAKULTÄT  
FÜR INFORMATIK  
Faculty of Informatics

Diplomarbeitspräsentation



# $k$ -LSM: A Relaxed Lock-Free Priority Queue

Masterstudium:  
Software Engineering & Internet Computing

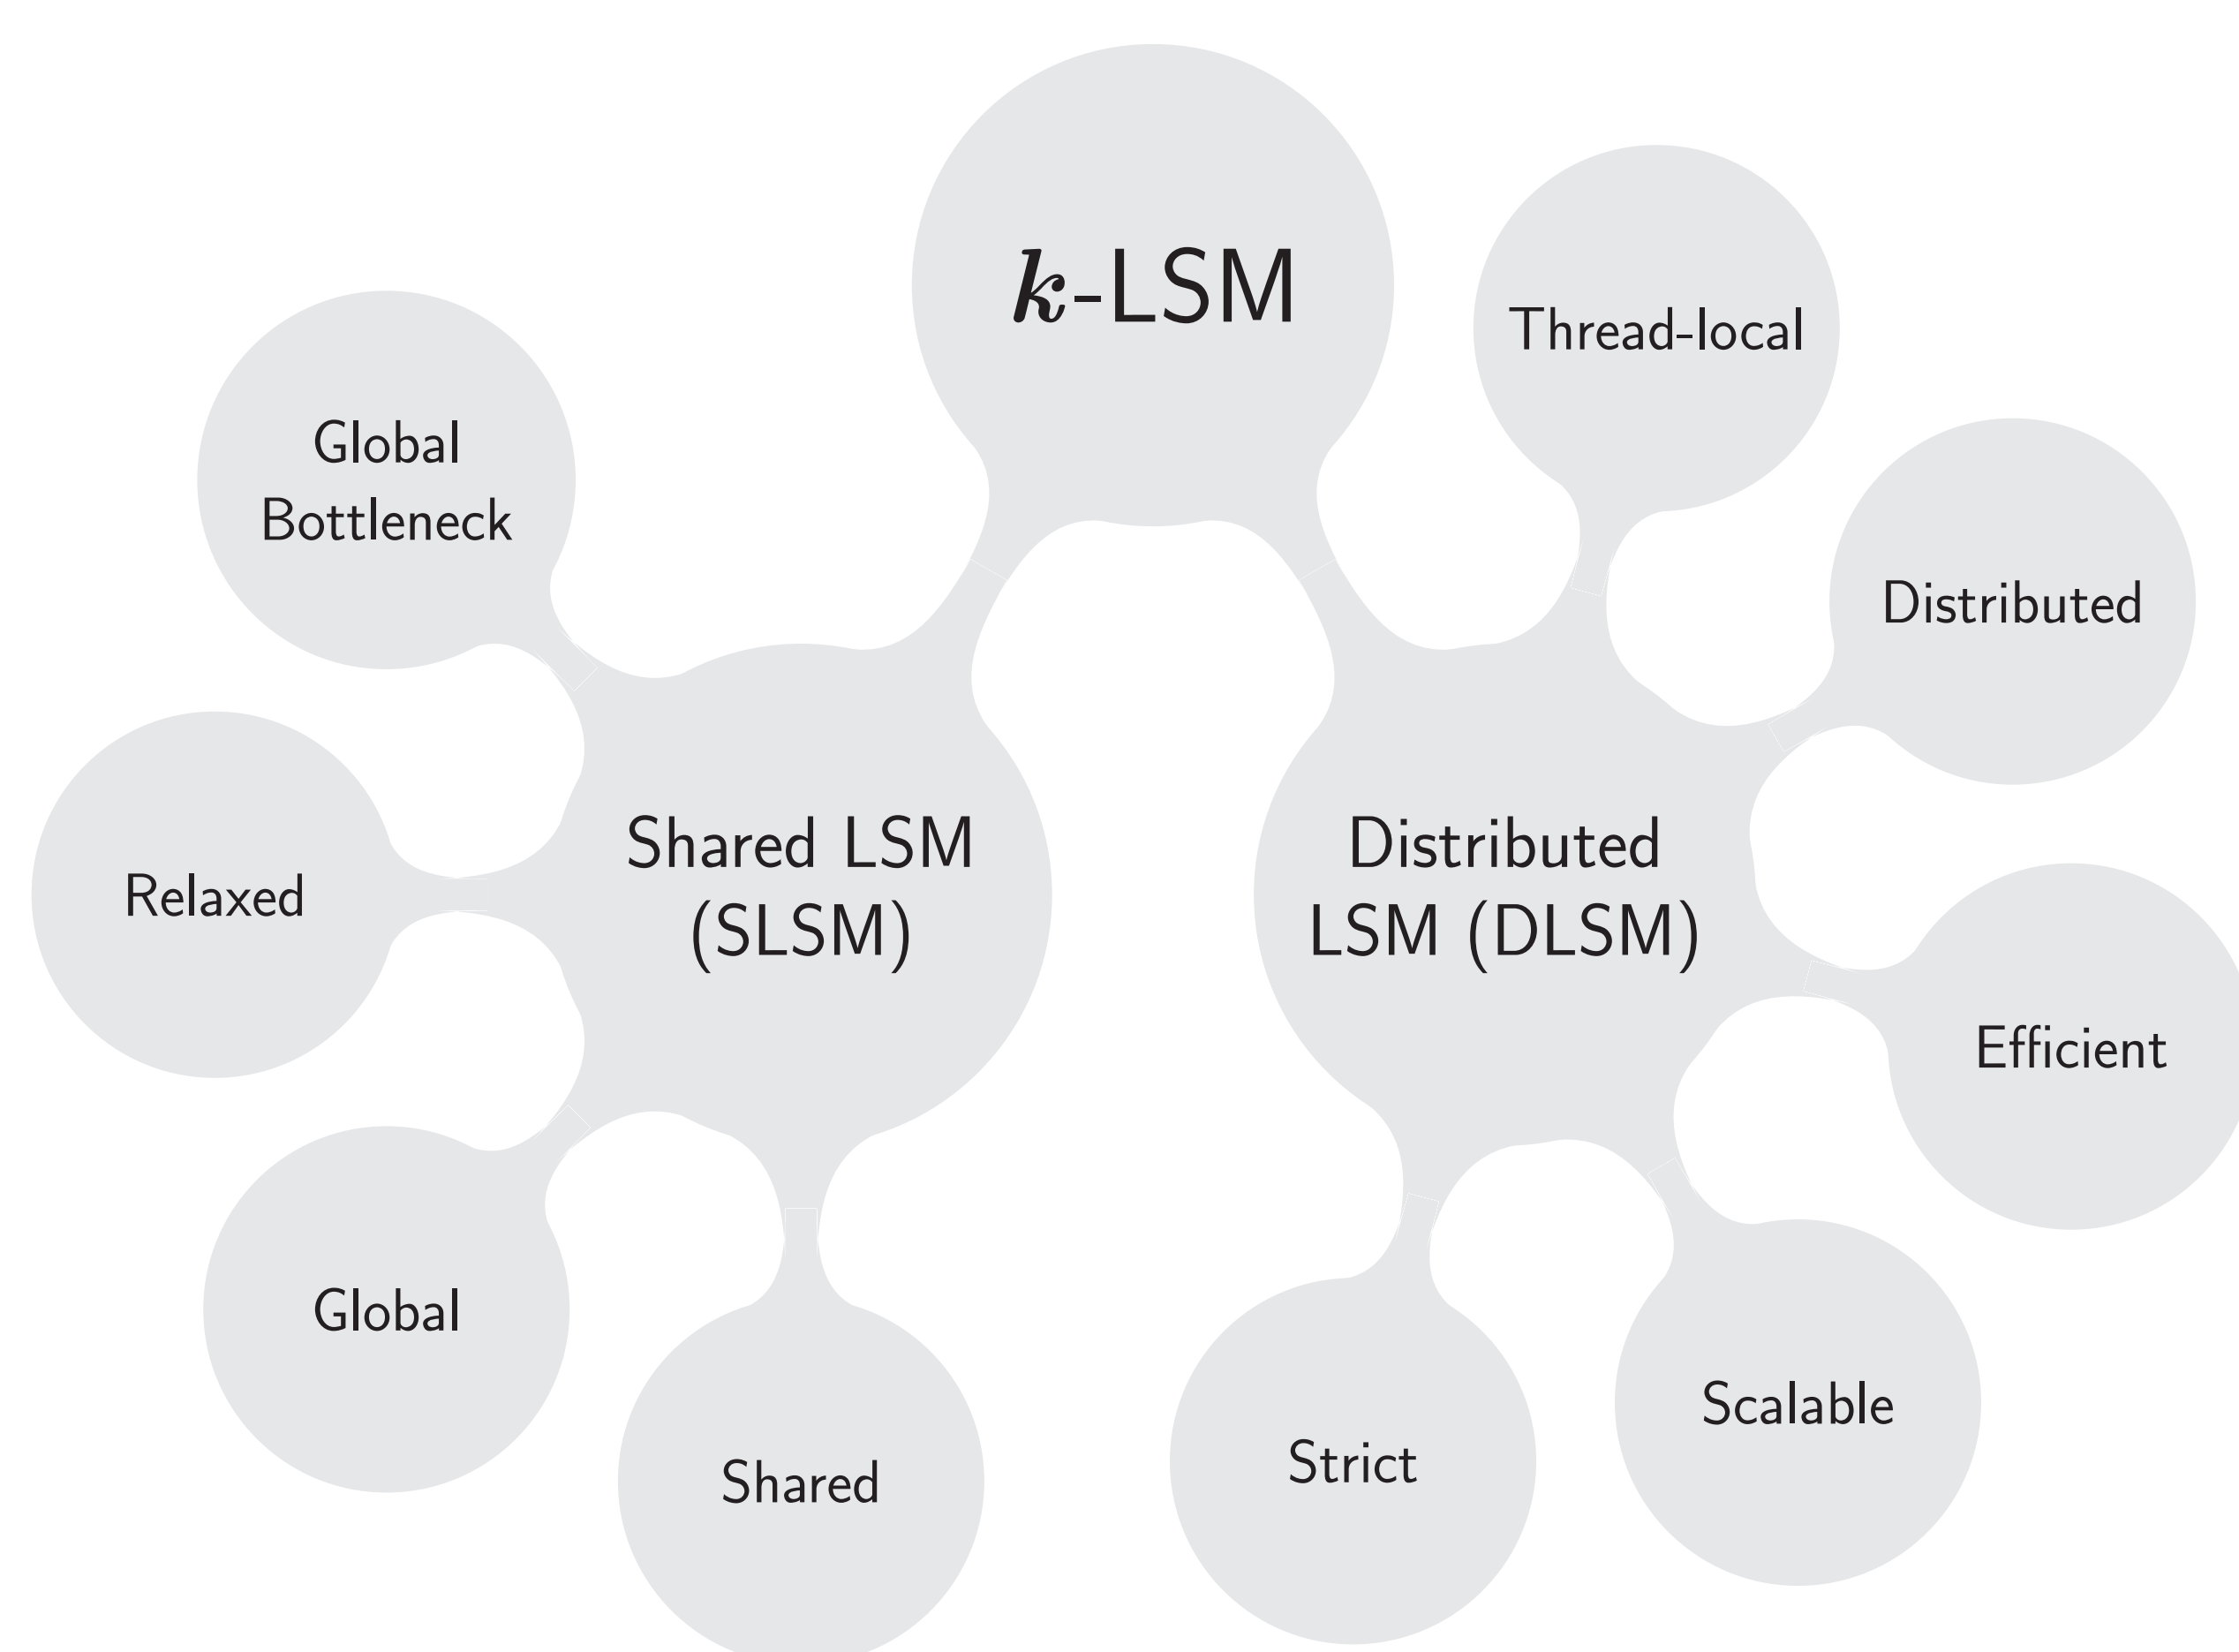
Jakob Gruber

Technische Universität Wien  
Institut für Informationssysteme  
Arbeitsbereich: Parallel Computing  
Betreuer: Prof. Dr. Scient. Jesper Larsson Träff

## Concurrent Priority Queues

- **Priority queues:** Abstract data structures for storing key/value pairs; efficient access to the item with the minimal (maximal) key.
- **Concurrent priority queues:** From early locking heaps, through lock-free SkipList-based designs, to current research into relaxed priority queues.
- **Lock-freedom:** Strong progress condition that can be efficient in practice.
- **Linearizability:** Each operation appears to take effect at some instant in time within the operation invocation.
- **Relaxed:** Weaken semantic guarantees in return for higher efficiency and/or scalability; useful for concurrent priority queues to alleviate the global bottleneck at the minimal (maximal) element.

## The $k$ -LSM Concurrent Priority Queue



The  $k$ -LSM (Log-structured Merge Tree) is a **relaxed**, **linearizable** and **lock-free** priority queue design by Wimmer, composed of two complementary priority queues; the SLSM, a centralized queue with strong guarantees but limited performance; and the DLSM, a distributed design without global guarantees but very high throughput.

- **Idea:** An LSM maintains items in a logarithmic number of ordered arrays to achieve amortized logarithmic times for deletions and insertions.
- **Relaxation:** Deletions may return one of the  $kP$  minimal items, where  $k$  is a configuration parameter and  $P$  is the number of threads.
- **Cache-efficient:** Usage of arrays and the standard merge algorithm result in good cache locality properties and high performance.

## Publications

- Martin Wimmer, Jakob Gruber, Jesper Larsson Träff, and Philippos Tsigas. “The lock-free  $k$ -LSM relaxed priority queue”. In: *Proceedings of the 20th ACM SIGPLAN Symposium on Principles and Practice of Parallel Programming*. ACM. 2015, pp. 277–278.

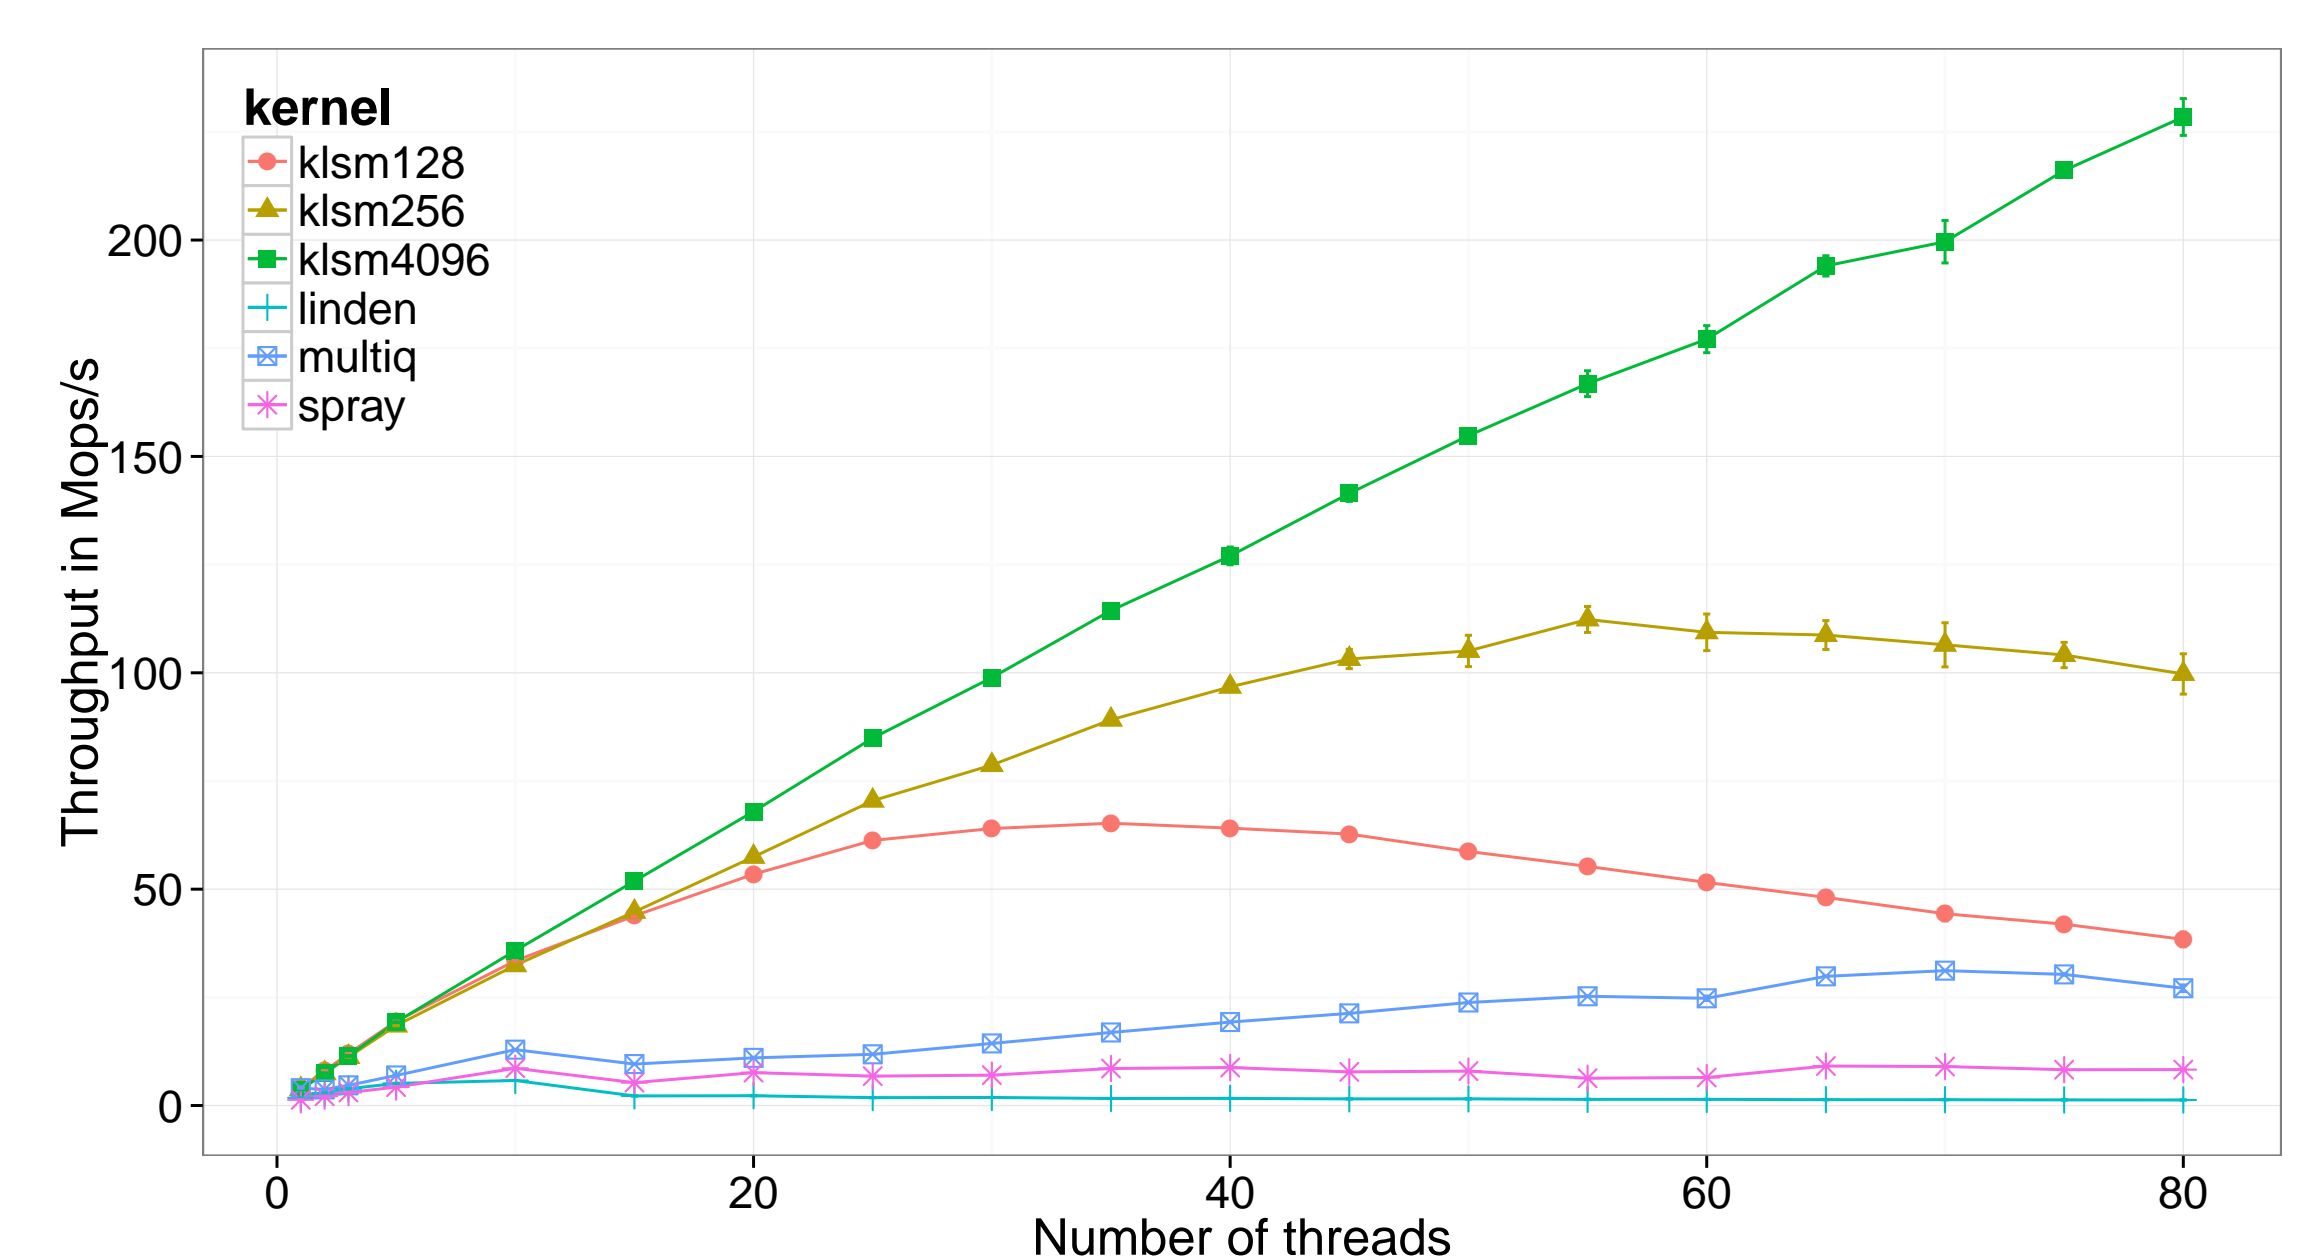
## Our Implementation

We improve on the  $k$ -LSM by providing a new, efficient, standalone implementation in addition to the previously available version as part of the task-scheduling framework *Pheet*. Our implementation is:

- **Scalable:** Under ideal circumstances and with sufficient relaxation, the  $k$ -LSM scales until the maximal thread-count on all tested machines.
- **Performant:** Outperforms the best other concurrent priority queues by up to a factor of ten.
- **Maintainable:** Uses classical software engineering concepts such as composability, readability, and separation of concerns.
- **Portable:** Implemented using C++11 atomics without reliance on external libraries.

## Evaluation & Results

- **Algorithms:** Evaluated against other state of the art designs, including the SprayList (relaxed, lock-free, probabilistic guarantees), the Lindén queue (strict, lock-free, linearizable), and Multiqueues (relaxed, lock-based, no guarantees).
- **Benchmarks:** Evaluation under a variety of scenarios, including balanced and unbalanced workloads combined with different key generation algorithms. We demonstrate that these parameters have a high impact on the performance of some data structures.
- **Throughput benchmarks** measure performance in millions of operations per second. Results are varied: in evenly distributed settings, the  $k$ -LSM outperforms all other data structures by a factor of 10, while in others, performance suffers since more stress is placed on the centralized SLSM.



- **Quality benchmarks** measure ranks of returned items, e.g., for each returned item, how many other items within the priority queue were skipped when determining it as the minimal item. The  $k$ -LSM more than fulfills its semantic guarantees, usually delivering results within the best fifth percentile.

Rank Error (in # items)	20 threads		40 threads		80 threads	
	Mean	StdDev	Mean	StdDev	Mean	StdDev
globallock	1.9	1.0	2.9	1.7	4.8	2.9
klsm128	33	31	55	46	430	294
klsm256	42	42	71	61	750	828
klsm4096	297	496	625	1014	10353	12667
multiq	984	2899	2252	7433	3787	12549