

k -LSM

A Relaxed Lock-Free Priority Queue

Jakob Gruber

February 26, 2016

Introduction

The k -LSM

- ▶ Concurrent priority queue by Wimmer
- ▶ *Lock-free*: Progress condition. At least one thread makes progress at any time.
- ▶ *Linearizable*: Safety condition. Operations appear to take effect at one instant in time.
- ▶ *Relaxed*: Deletions may return one of kP smallest elements (k is configurable, P is the thread count).

Introduction

Motivation




Figure: Throughput on mars, uniform workload, uniform key generation.

Introduction

Motivation




Figure: Throughput on mars, uniform workload, uniform key generation.

Introduction

Motivation




Figure: Throughput on mars, uniform workload, uniform key generation.

Introduction

Motivation




Figure: Throughput on mars, uniform workload, uniform key generation.

Introduction

In this thesis

Previously, k -LSM only part of task-scheduling framework *Pheet*.

- ▶ Standalone reimplementation (C++11, 4800 SLOC).
- ▶ Extensive evaluation against related structures on multiple machines.
- ▶ Different benchmarks demonstrate behavioral changes.
- ▶ Semantic quality benchmarks show results are within theoretical bounds.

Introduction

Contents

- ▶ Background.
- ▶ Related work.
- ▶ The k -LSM design.
- ▶ Results.

Introduction

Sequential priority queues

Priority Queues (PQs):

- ▶ Standard abstract data structure.
- ▶ Interface consists of two $O(\log n)$ operations:

```
void insert(const K &key, const V &val);  
bool delete_min(V &val);
```

- ▶ Typical backing data structures: heaps & search trees.

Related Work

- ▶ Strict concurrent PQs
 - ▶ Hunt et al.: Fine-grained locking heap.
 - ▶ Shavit and Lotan: First SkipList-based PQ.
 - ▶ Sundell and Tsigas: First lock-free PQ.
 - ▶ Lindén and Jonsson: Minimizes contention.
 - ▶ Braginsky: List of arrays, efficient.
- ▶ Relaxed concurrent PQs
 - ▶ Alistarh et al.: SprayList.
 - ▶ Rihani, Sanders, and Dementiev: Multiqueues.
 - ▶ Wimmer et al.: k-Log-structured Merge Tree (LSM).

Relaxed Priority Queues

- ▶ Strict PQs have inherent bottleneck at minimal element.
- ▶ Another approach is to relax semantics, i.e. instead of returning *the* minimal element, return one of the k minimal elements.

Relaxed Priority Queues

Alistarh et al.: SprayList




Figure: The SprayList is based on a lock-free SkipList. Deletions randomly pick one of the $O(P \log^3 P)$ smallest items.

Relaxed Priority Queues

Rihani, Sanders, and Dementiev: Multiqueues


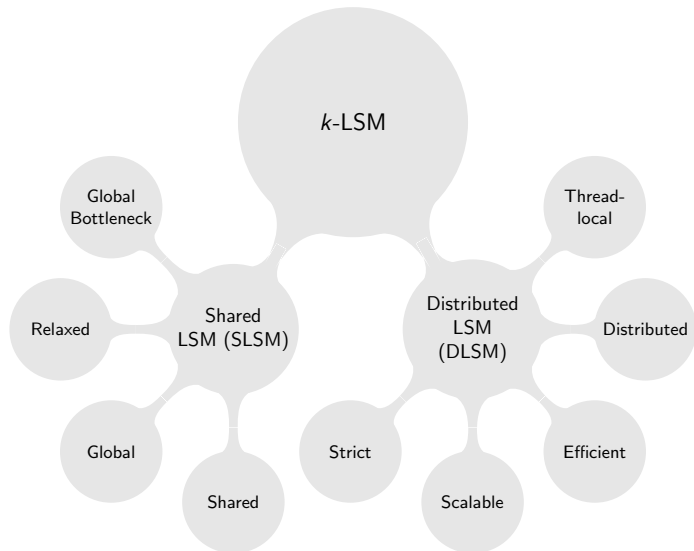


Figure: Multiqueues are a collection of priority queues. Deletions remove the minimal item from a randomly selected queue.

k -LSM




k -LSM

The LSM


- ▶ Based on the concept of LSMs.
- ▶ Maintain items in a logarithmic number of sorted arrays.
- ▶ Array merges are the central operation.
- ▶ Both insertions and deletions are in $O(\log n)$.

k -LSM

LSM Insertions



(a) The initial state.




(b) Item 7 inserted as a new block.


Figure: Insertion of a new element into the LSM.

k -LSM

LSM Insertions



(a) After the first merge.




(b) After the second merge.


Figure: Insertion of a new element into the LSM.

k -LSM

LSM Deletions



(a) The initial state.



(b) Item 0 marked as deleted.

Figure: Deletion from the LSM.

k -LSM

DLSM

- ▶ One LSM per thread (thread-local).
- ▶ Inter-thread communication only when local LSM is empty (spy).

k -LSM

SLSM

- ▶ One global LSM shared by all threads.
- ▶ Relaxation through so-called *pivot* ranges.

k -LSM

SLSM




Figure: Pivot ranges contain subset of $k + 1$ smallest elements.

k -LSM

k -LSM: The kP bound

- ▶ k -LSM deletions peek at DLSM and SLSM and remove the smaller item.
- ▶ Each thread-local DLSM has a capacity of $k \rightarrow$ deletions skip at most $k(P - 1)$ items.
- ▶ The SLSM pivot range contains $k + 1$ smallest items \rightarrow deletions skip at most k items.
- ▶ Combining both: at most kP items skipped.

Results

Throughput: Uniform workload, uniform key generation

- ▶ Most commonly used benchmark.
- ▶ Throughput: Number of operations per second.
- ▶ Each thread: 50% insertions, 50% deletions.
- ▶ Keys taken at random from 32-bit integer range.

Results

Throughput: Uniform workload, uniform key generation




Figure: mars (80 core Intel Xeon at 2 GHz)

Results

Throughput: Uniform workload, ascending keygeneration

- ▶ Keys taken from $[0, 512[+ t$.
- ▶ Induces FIFO-like behavior.

Results

Throughput: Uniform workload, ascending keygeneration




Figure: mars.

Results

Throughput: Split workload, uniform key generation

- ▶ 50% of threads insert, others delete.

Results

Throughput: Split workload, uniform key generation




Figure: mars.

Results

Throughput: Split workload, ascending keygeneration




Figure: mars.

Results

- ▶ Performance directly dependent on utilization of the DLSSM.
- ▶ Uniform/uniform benchmark: 95% of all deletions from the DLSSM.

Results

Quality

- ▶ And how good is the quality of results?
- ▶ Determined by rank: for each `delete_min`, if the removed item is the k -smallest item at that time, it has rank k .
- ▶ k -LSM results are usually within 5th percentile of allowed relaxation.
- ▶ `multiq` approximately comparable to `k1sm4096`.

Conclusion

- ▶ Standalone reimplementation completed successfully.
- ▶ The k -LSM can have exceptional performance.
- ▶ ... under the right conditions.
- ▶ The standard uniform workload, uniform key generation benchmark is not enough to evaluate performance.
- ▶ Very good quality results in practice.
- ▶ Multiqueues have lower peak performance, but have stable behavior in all experiments.

Conclusion

Questions?

Results

Throughput: Uniform workload, uniform key generation




Figure: mars, comparison against the Pheet k -LSM

Results

Throughput: Uniform workload, uniform key generation




Figure: saturn

Results

Throughput: Uniform workload, uniform key generation




Figure: ceres

Results

Quality: Uniform workload, uniform key generation

	20 threads		40 threads		80 threads	
	Mean	SD	Mean	SD	Mean	SD
klsm256	42	42	68	57	635	464
klsm4096	422	729	1124	1287	13469	13980
multiq	1163	3607	2296	7881	3753	12856

Table: mars.